

Approximate Principal Components Analysis of Large Data Sets

Daniel J. McDonald

Department of Statistics
Indiana University

<http://mypage.iu.edu/~dajmcdon>
dajmcdon@indiana.edu

Joint work with:
Darren Homrighausen

August 5, 2014

DATA REDUCTION

Modern statistical applications — genomics, neural image analysis, text analysis — have large numbers of covariates p

Also frequently have lots of observations n .

Need to do some dimension reduction

Many methods — multidimensional scaling, discriminant analysis, locally linear embeddings, Laplacian eigenmaps, and many others

NOTATION

- Observe $\tilde{X}_1, \dots, \tilde{X}_n$, where $\tilde{X}_i \in \mathbb{R}^p$
- Form $\tilde{\mathbf{X}} = [\tilde{X}_1 \ \dots \ \tilde{X}_n]^\top \in \mathbb{R}^{n \times p}$
- Call $\mathbf{X} = \tilde{\mathbf{X}} - \bar{\mathbf{X}}$ where $\bar{\mathbf{X}} = n^{-1} \mathbf{1} \mathbf{1}^\top \tilde{\mathbf{X}}$
- Write the SVD of \mathbf{X} as

$$\mathbf{X} = \mathbf{U} \mathbf{\Lambda} \mathbf{V}^\top,$$

- For general rank r matrices \mathbf{A} we write

$$\mathbf{A} = U(\mathbf{A}) \mathbf{\Lambda}(\mathbf{A}) V(\mathbf{A})^\top \quad \text{where} \quad \mathbf{\Lambda}(\mathbf{A}) = \text{diag}(\lambda_1(\mathbf{A}), \dots, \lambda_r(\mathbf{A}))$$

- For some matrix \mathbf{A} , we use \mathbf{A}_d to be the first d columns of \mathbf{A} .

PRINCIPAL COMPONENTS ANALYSIS

- For $d \leq r$, PCA gives a projection that minimizes the squared error distance between the data and the projection
- Projections involve \mathbf{U}_d , \mathbf{V}_d , and Λ_d

EXAMPLE

PCA for regression with response vector Y , the estimated coefficients would be

$$\hat{\beta}_{pcr} = \mathbf{V}_d \Lambda_d^\dagger \mathbf{U}_d^\top Y = \mathbf{V}_d \mathbf{V}_d^\top \mathbf{X}^\dagger Y \quad (1)$$

and the fitted values can be written $\hat{Y} = \mathbf{U}_d \mathbf{U}_d^\top Y$

A BIG PROBLEM

- PCA requires computing an SVD
- So do most other data reduction methods
- Datasets are large: a recent ImageNet contest had $n \approx 10^6$ and $p \approx 65000$ ¹
- That was only about 8% of the data set.
- SVD requires $O(np^2 + n^3)$ computations **check me**
- And you need to store the entire matrix in fast memory
- This is bad.

¹ see e.g. Krizhevsky, Sutskever, and Hinton. *NIPS* 2013

A BIG PROBLEM

- PCA requires computing an SVD
- So do most other data reduction methods
- Datasets are large: a recent ImageNet contest had $n \approx 10^6$ and $p \approx 65000$ ¹
- That was only about 8% of the data set.
- SVD requires $O(np^2 + n^3)$ computations **check me**
- And you need to store the entire matrix in fast memory
- **This is bad.**

¹ see e.g. Krizhevsky, Sutskever, and Hinton. *NIPS* 2013

APPROXIMATIONS

Can't take the SVD of \mathbf{X}

What about approximating it?

Focus on two methods of “approximate SVD”

- 1 Nyström extension
- 2 Column sampling

Not our methods.

A QUICK “SKETCH” OF THE INTUITION

- Both methods fall into a larger class of methods
- Suppose we want to approximate $\mathbf{A} \in \mathbb{R}^{q \times q}$
- Assume \mathbf{A} is symmetric and positive semi-definite
- Choose $l \ll q$ and form a “sketching” matrix $\Phi \in \mathbb{R}^{q \times l}$
- Then write $\mathbf{A} \approx (\mathbf{A}\Phi)(\Phi^\top \mathbf{A}\Phi)^\dagger (\mathbf{A}\Phi)^\top$.
- Different Φ yield different approximations
- For Nyström and column sampling use

$$\Phi = \pi \tau$$

Where π is a permutation of \mathbf{I}_q and $\tau = [\mathbf{I}_l \quad \mathbf{0}]^\top$.

NOTES ON SKETCHING

- We don't actually do this, but this is the generality
- Essentially, let

$$\mathbf{S} = \frac{1}{n} \mathbf{X}^\top \mathbf{X} \quad \text{and} \quad \mathbf{Q} = \mathbf{X} \mathbf{X}^\top$$

- Both of these are symmetric, positive semi-definite
- Randomly choose l entries in $\{1, \dots, n\}$ and $\{1, \dots, p\}$
- Then partition the matrix so the selected portion is \mathbf{S}_{11} and \mathbf{Q}_{11}

$$\mathbf{S} = \mathbf{V} \Lambda^2 \mathbf{V}^\top = \begin{bmatrix} \mathbf{S}_{11} & \mathbf{S}_{12} \\ \mathbf{S}_{21} & \mathbf{S}_{22} \end{bmatrix} \quad \mathbf{Q} = \mathbf{U} \Lambda^2 \mathbf{U}^\top = \begin{bmatrix} \mathbf{Q}_{11} & \mathbf{Q}_{12} \\ \mathbf{Q}_{21} & \mathbf{Q}_{22} \end{bmatrix}$$

APPROXIMATING THINGS WE DON'T CARE ABOUT

If we want to approximate \mathbf{S} (or \mathbf{Q}), we have for example

Nyström

$$\mathbf{S} \approx \begin{bmatrix} \mathbf{S}_{11} \\ \mathbf{S}_{21} \end{bmatrix} \mathbf{S}_{11}^\dagger \begin{bmatrix} \mathbf{S}_{11} & \mathbf{S}_{12} \end{bmatrix}$$

Column sampling

$$\mathbf{S} \approx U \left(\begin{bmatrix} \mathbf{S}_{11} \\ \mathbf{S}_{21} \end{bmatrix} \right) \Lambda \left(\begin{bmatrix} \mathbf{S}_{11} \\ \mathbf{S}_{21} \end{bmatrix} \right) U \left(\begin{bmatrix} \mathbf{S}_{11} \\ \mathbf{S}_{21} \end{bmatrix} \right)^\top$$

Previous theoretical results have focused on the accuracy of these approximations and the way to randomly select the indices

WHAT WE ACTUALLY WANT...

We really want \mathbf{U} , \mathbf{V} , and Λ . Then we can get the principal components, principal coordinates, and the amount of variance explained.

It turns out that there are quite a few ways to use these two methods to get the things we want.

Let

$$L(\mathbf{S}) = \begin{bmatrix} \mathbf{S}_{11} \\ \mathbf{S}_{21} \end{bmatrix}$$

$$L(\mathbf{Q}) = \begin{bmatrix} \mathbf{Q}_{11} \\ \mathbf{Q}_{21} \end{bmatrix}$$

LOTS OF APPROXIMATIONS

After some reasonable algebra...

Quantity of interest	Label	Approximations
\mathbf{V}	\mathbf{V}_{nys} \mathbf{V}_{cs}	$L(\mathbf{S})\mathbf{V}(\mathbf{S}_{11})\Lambda(\mathbf{S}_{11})^\dagger$ $\mathbf{U}(L(\mathbf{S}))$
\mathbf{U}	\mathbf{U}_{nys} \mathbf{U}_{cs} $\hat{\mathbf{U}}_{nys}$ $\hat{\mathbf{U}}_{cs}$ $\hat{\mathbf{U}}$	$L(\mathbf{Q})\mathbf{V}(\mathbf{Q}_{11})\Lambda(\mathbf{Q}_{11})^\dagger$ $\mathbf{U}(L(\mathbf{Q}))$ $\mathbf{X}\mathbf{V}_{nys}\Lambda_{nys}^{\dagger/2}$ $\mathbf{X}\mathbf{V}_{cs}\Lambda_{cs}^{\dagger/2}$ $\mathbf{U}(\mathbf{x}_1)$

LOTS OF APPROXIMATIONS

After some reasonable algebra...

Quantity of interest	Label	Approximations
\mathbf{V}	\mathbf{V}_{nys} \mathbf{V}_{cs}	$L(\mathbf{S})\mathbf{V}(\mathbf{S}_{11})\Lambda(\mathbf{S}_{11})^\dagger$ $\mathbf{U}(L(\mathbf{S}))$
\mathbf{U}	\mathbf{U}_{nys} \mathbf{U}_{cs} $\hat{\mathbf{U}}_{nys}$ $\hat{\mathbf{U}}_{cs}$ $\hat{\mathbf{U}}$	$L(\mathbf{Q})\mathbf{V}(\mathbf{Q}_{11})\Lambda(\mathbf{Q}_{11})^\dagger$ $\mathbf{U}(L(\mathbf{Q}))$ $\mathbf{X}\mathbf{V}_{nys}\Lambda_{nys}^{\dagger/2}$ $\mathbf{X}\mathbf{V}_{cs}\Lambda_{cs}^{\dagger/2}$ $\mathbf{U}(\mathbf{x}_1)$

LOTS OF APPROXIMATIONS

After some reasonable algebra...

Quantity of interest	Label	Approximations
\mathbf{V}	\mathbf{V}_{nys} \mathbf{V}_{cs}	$L(\mathbf{S})\mathbf{V}(\mathbf{S}_{11})\Lambda(\mathbf{S}_{11})^\dagger$ $\mathbf{U}(L(\mathbf{S}))$
\mathbf{U}	\mathbf{U}_{nys} \mathbf{U}_{cs} $\hat{\mathbf{U}}_{nys}$ $\hat{\mathbf{U}}_{cs}$ $\hat{\mathbf{U}}$	$L(\mathbf{Q})\mathbf{V}(\mathbf{Q}_{11})\Lambda(\mathbf{Q}_{11})^\dagger$ $\mathbf{U}(L(\mathbf{Q}))$ $\mathbf{X}\mathbf{V}_{nys}\Lambda_{nys}^{\dagger/2}$ $\mathbf{X}\mathbf{V}_{cs}\Lambda_{cs}^{\dagger/2}$ $\mathbf{U}(\mathbf{x}_1)$

LOTS OF APPROXIMATIONS

After some reasonable algebra...

Quantity of interest	Label	Approximations
\mathbf{V}	\mathbf{V}_{nys} \mathbf{V}_{cs}	$L(\mathbf{S})\mathbf{V}(\mathbf{S}_{11})\Lambda(\mathbf{S}_{11})^\dagger$ $\mathbf{U}(L(\mathbf{S}))$
\mathbf{U}	\mathbf{U}_{nys} \mathbf{U}_{cs} $\hat{\mathbf{U}}_{nys}$ $\hat{\mathbf{U}}_{cs}$ $\hat{\mathbf{U}}$	$L(\mathbf{Q})\mathbf{V}(\mathbf{Q}_{11})\Lambda(\mathbf{Q}_{11})^\dagger$ $\mathbf{U}(L(\mathbf{Q}))$ $\mathbf{X}\mathbf{V}_{nys}\Lambda_{nys}^{\dagger/2}$ $\mathbf{X}\mathbf{V}_{cs}\Lambda_{cs}^{\dagger/2}$ $\mathbf{U}(\mathbf{x}_1)$

LOTS OF APPROXIMATIONS

After some reasonable algebra...

Quantity of interest	Label	Approximations
\mathbf{V}	\mathbf{V}_{nys} \mathbf{V}_{cs}	$L(\mathbf{S})\mathbf{V}(\mathbf{S}_{11})\Lambda(\mathbf{S}_{11})^\dagger$ $\mathbf{U}(L(\mathbf{S}))$
\mathbf{U}	\mathbf{U}_{nys} \mathbf{U}_{cs} $\hat{\mathbf{U}}_{nys}$ $\hat{\mathbf{U}}_{cs}$ $\hat{\mathbf{U}}$	$L(\mathbf{Q})\mathbf{V}(\mathbf{Q}_{11})\Lambda(\mathbf{Q}_{11})^\dagger$ $\mathbf{U}(L(\mathbf{Q}))$ $\mathbf{X}\mathbf{V}_{nys}\Lambda_{nys}^{\dagger/2}$ $\mathbf{X}\mathbf{V}_{cs}\Lambda_{cs}^{\dagger/2}$ $\mathbf{U}(\mathbf{x}_1)$



Method	Complexity:	
	Computational	Storage
Nyström	$O(nl^2 + l^3)$	$O(l^2)$ [$O(nl)$]
Column sampling	$O(lnp + p^2l)$	$O(pl)$

- Column sampling results in orthogonal \mathbf{U} and \mathbf{V} , Nyström doesn't
- Using \mathbf{Q} is weird. Thus the “hatted \mathbf{U} ” versions.
- $\hat{\mathbf{U}} = U(\mathbf{x}_1)$ where $\mathbf{X} = [\mathbf{x}_1 \quad \mathbf{x}_2]$ seems reasonable if we think that only the first l selected covariates matter (gets used in supervised PCA)²

² Bair, Hastie, Paul, and Tibshirani, *JASA*, 2006.

SO WHAT DO WE USE?

Well...

It depends. We did some theory which gives a way to calculate how far off your approximation might be. You could use these bounds if you like to use your data and make a choice.

Did some simulations too.


- Draw $\tilde{X}_i \stackrel{iid}{\sim} N_p(0, \Sigma)$ for $n = 5000$, $p = 3000$.
- 4 conditions: **Random**_{0.001}, **Random**_{0.01}, **Random**_{0.1}, and **Band**.
 - **Random** _{x} — with probability x and for $i < j$, $\Sigma_{ij}^{-1} = 1$ and 0 otherwise. Diagonal elements are always equal to 1. And symmetrize.
 - **Band** — $\Sigma_{ij}^{-1} = 1$ if $|i - j| \leq 50$ and 0 otherwise.
- The graphical models have approximately $\binom{p}{2} \cdot x$ edges for **Random** _{x} and exactly $25(2p - 1 - 50)$ edges for **Band**.

WHAT WE REPORT


In the \mathbf{V} case

$$\frac{\|\mathbf{V}_{nys,d}(\mathbf{V}_{nys,d}^\top \mathbf{V}_{nys,d})^{-1} \mathbf{V}_{nys,d}^\top - \mathbf{V}_d \mathbf{V}_d^\top\|_F}{\|\mathbf{V}_{cs,d} \mathbf{V}_{cs,d}^\top - \mathbf{V}_d \mathbf{V}_d^\top\|_F}.$$


APPROXIMATIONS TO V




($d = 2$)



($d = 5$)




($d = 10$)



($d = 15$)

- y -axis, relative performance of Nyström to column sampling.
- Less than 1—better. More than 1—worse.
- x -axis, approximation parameter l (grid based on d).
- Random_{0.001} (solid, square), Random_{0.01} (dashed, circle), Random_{0.1} (dotted, triangle), and Band(dot-dash, diamond).
- d small, computation time is similar. d large, Nyström is 10–15% of column sampling

APPROXIMATIONS TO \mathbf{U} ($d = 10$)



- y -axis, performance relative to $\hat{\mathbf{U}}_{cs}$
- x -axis, approximation parameter l
- $\hat{\mathbf{U}}$, \mathbf{U}_{nys} , $\hat{\mathbf{U}}_{nys}$, $\hat{\mathbf{U}}_{cs}$
- Clockwise from top left: Random_{0.001}, Random_{0.01}, Band, Random_{0.1}

CONCLUSIONS

- For computing \mathbf{V} , CS beats Nyström in terms of accuracy, but is much slower for similar choices of the approximation parameter and d large.
- For computing \mathbf{U} , the naïve methods are bad, better to approximate \mathbf{V} and multiply, so see above
- $\hat{\mathbf{U}}$ really stinks. This is used for supervised PCA. Future research (don't steal, this one's ours)
- Other choices of Φ (also future research)
- Thanks NSF for funding us.
- For more info, see the paper. Contains boring theory, more extensive simulations, and application to Enron email dataset. (resubmitted to JCGS)